

Parallel Computing B-PB20000178

李笑

Lab 1 - OpenMP及CUDA实验环境的搭建

以下两张截图是我个人电脑的配置：

TechPowerUp GPU-Z 2.52.0


Graphics Card | Sensors | Advanced | Validation

Name: Intel(R) UHD Graphics [Lookup](#)

GPU: Comet Lake GT2 Revision: V0

Technology: 14 nm Die Size: Unknown

Release Date: Aug 21, 2019 Transistors: Unknown

BIOS Version: Unknown  ☒ UEFI

Subvendor: Huaqin Device ID: 8086 9B41 - 1E83 3E1B

ROPs/TMUs: 8 / 16 Bus Interface: N/A ?

Shaders: 24 Unified DirectX Support: 12 (12_1)

Pixel Fillrate: 9.2 GPixel/s Texture Fillrate: 18.4 GTexel/s

Memory Type: LPDDR3 Bus Width: 128 bit

Memory Size: N/A Bandwidth: 34.0 GB/s

Driver Version: 27.20.100.8984 DCH / Win10 64

Driver Date: Nov 19, 2020 Digital Signature: WHQL


GPU Clock: 299 MHz Memory: 1064 MHz Boost: 1147 MHz

Default Clock: 300 MHz Memory: 1067 MHz Boost: 1150 MHz

Multi-GPU: Disabled Resizable BAR: Disabled

Computing ☒ OpenCL ☐ CUDA ☒ DirectCompute ☒ DirectML

Technologies ☒ Vulkan ☐ Ray Tracing ☒ PhysX ☒ OpenGL 4.6

Intel(R) UHD Graphics  [Close](#)

TechPowerUp GPU-Z 2.52.0


Graphics Card | Sensors | Advanced | Validation

Name: NVIDIA GeForce MX350 [Lookup](#)

GPU: GP107 Revision: A1

Technology: 14 nm Die Size: 132 mm²

Release Date: Feb 10, 2020 Transistors: 3300M

BIOS Version: 86.07.92.00.7A  ☐ UEFI

Subvendor: Huaqin Device ID: 10DE 1C94 - 1E83 3E1B

ROPs/TMUs: 16 / 40 Bus Interface: PCIe x16 3.0 @ x4 1.1 ?

Shaders: 640 Unified DirectX Support: 12 (12_1)

Pixel Fillrate: 23.5 GPixel/s Texture Fillrate: 58.7 GTexel/s

Memory Type: GDDR5 (Hynix) Bus Width: 64 bit

Memory Size: 2048 MB Bandwidth: 56.1 GB/s

Driver Version: 26.21.14.4250 (NVIDIA 442.50) DCH / Win10 64

Driver Date: Feb 24, 2020 Digital Signature: WHQL


GPU Clock: 1354 MHz Memory: 1752 MHz Boost: 1468 MHz

Default Clock: 1354 MHz Memory: 1752 MHz Boost: 1468 MHz

NVIDIA SLI: Disabled Resizable BAR: Disabled

Computing: ☒ OpenCL ☒ CUDA ☒ DirectCompute ☒ DirectML

Technologies: ☒ Vulkan ☐ Ray Tracing ☒ PhysX ☒ OpenGL 4.6

NVIDIA GeForce MX350  [Close](#)

切换到Ubuntu20.04进行后续实验，通过命令行查看配置：

```
xiaoli@xiaoli-KLVC-WXX9: ~  
xiaoli@xiaoli-KLVC-WXX9:~$ cat /proc/cpuinfo | grep "physical id" | sort | uniq | wc -l  
1  
xiaoli@xiaoli-KLVC-WXX9:~$ cat /proc/cpuinfo | grep "cpu cores" | uniq  
cpu cores      : 4  
xiaoli@xiaoli-KLVC-WXX9:~$ cat /proc/cpuinfo | grep "processor" | wc -l  
8  
xiaoli@xiaoli-KLVC-WXX9:~$ cat /proc/cpuinfo | grep name | cut -f2 -d: | uniq -c  
      8 Intel(R) Core(TM) i7-10510U CPU @ 1.80GHz
```

安装OpenMP

步骤：

1. 快捷键Ctrl+Atl+T打开终端
2. 在终端输入`sudo apt-get install libomp-dev`安装OpenMP
3. 在终端输入`sudo apt-get install gcc`安装GCC
4. 在终端输入`gcc --version`检查安装是否成功

```
xiaoli@xiaoli-KLVC-WXX9: ~  
xiaoli@xiaoli-KLVC-WXX9:~$ gcc --version  
gcc (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0  
Copyright (C) 2019 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions. There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

5. 在终端输入`echo |cpp -fopenmp -dM |grep -i open`检查OpenMP安装是否成功

```
xiaoli@xiaoli-KLVC-WXX9: ~  
xiaoli@xiaoli-KLVC-WXX9:~$ echo |cpp -fopenmp -dM |grep -i open  
#define _OPENMP 201511
```

安装CUDA

配置前：

```
xiaoli@xiaoli-KLVC-WXX9: ~  
xiaoli@xiaoli-KLVC-WXX9:~$ cat /usr/local/cuda/version.txt  
cat: /usr/local/cuda/version.txt: No such file or directory  
xiaoli@xiaoli-KLVC-WXX9:~$ lspci | grep -i nvidia  
01:00.0 3D controller: NVIDIA Corporation GP107M [GeForce MX350] (rev a1)  
xiaoli@xiaoli-KLVC-WXX9:~$ nvidia-smi  
  
Command 'nvidia-smi' not found, but can be installed with:  
  
sudo apt install nvidia-340 # version 340.108-0ubuntu5.20.04.2, or  
sudo apt install nvidia-utils-390 # version 390.157-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-450-server # version 450.216.04-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-470 # version 470.161.03-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-470-server # version 470.161.03-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-510 # version 510.108.03-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-515 # version 515.86.01-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-515-server # version 515.86.01-0ubuntu0.20.04.3  
sudo apt install nvidia-utils-525 # version 525.89.02-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-525-server # version 525.85.12-0ubuntu0.20.04.1  
sudo apt install nvidia-utils-435 # version 435.21-0ubuntu7  
sudo apt install nvidia-utils-440 # version 440.82+really.440.64-0ubuntu6  
sudo apt install nvidia-utils-418-server # version 418.226.00-0ubuntu0.20.04.2
```

上图信息表明，我的电脑装有NVIDIA显卡，但是没有安装显卡驱动

步骤：

1. 手动安装显卡驱动。依次在终端输入，选择系统推荐版本驱动 `nvidia-driver-525`

```
$ sudo add-apt-repository ppa:graphics-drivers/ppa  
$ sudo apt update  
$ ubuntu-drivers devices  
$ sudo apt install nvidia-driver-525
```

```

+ xiaoli@xiaoli-KLVC-WXX9: ~
+ xiaoli@xiaoli-KLVC-WXX9:~$ nvidia-smi
Tue Apr  4 14:55:40 2023
+-----+
| NVIDIA-SMI 525.105.17    Driver Version: 525.105.17    CUDA Version: 12.0    |
+-----+-----+-----+-----+-----+
| GPU  Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                               |                    |    MIG M. |
+-----+-----+-----+-----+-----+
|   0  NVIDIA GeForce ...  Off        | 00000000:01:00.0 Off |           N/A       |
| N/A   37C    P8      N/A /  N/A |  9MiB / 2048MiB |      0%      Default |
|                               |                    |           N/A       |
+-----+-----+-----+-----+-----+

Processes:
+-----+-----+-----+-----+-----+
| GPU  GI   CI           PID  Type  Process name                      GPU Memory |
|      ID  ID                                     Usage      |
+-----+-----+-----+-----+-----+
|   0  N/A  N/A       1158    G    /usr/lib/xorg/Xorg                 4MiB      |
|   0  N/A  N/A       1839    G    /usr/lib/xorg/Xorg                 4MiB      |
+-----+-----+-----+-----+-----+

```

2. 关闭系统自带驱动nouveau。通过在终端输入指令`lsmod | grep nouveau`查看驱动启用情况。我输入后发现有输出，表明nouveau驱动正在工作。所以,接下来在终端输入`sudo gedit /etc/modprobe.d/blacklist.conf`，弹出了blacklist.conf文件，在文件末尾加上blacklist

nouveau和options nouveau modeset=0两行并保存。

```

+ xiaoli@xiaoli-KLVC-WXX9: ~
xiaoli@xiaoli-KLVC-WXX9:~$ lsmod | grep nouveau
nouveau                2285568    1
mxm_wmi                  16384     1 nouveau
drm_ttm_helper           16384     1 nouveau
ttm                      86016     3 drm_ttm_helper,i915,nouveau
drm_kms_helper           307200    2 i915,nouveau
i2c_algo_bit             16384     2 i915,nouveau
drm                     618496   20 drm_kms_helper,drm_ttm_helper,i915,ttm,nouveau
video                    57344     3 int3406_thermal,i915,nouveau
wmi                      32768     5 intel_wmi_thunderbolt,huawei_wmi,wmi_bmf,mxm_wmi,nouveau
nouveau
xiaoli@xiaoli-KLVC-WXX9:~$ sudo gedit /etc/modprobe.d/blacklist.conf
[sudo] password for xiaoli:

(gedit:10075): Tepl-WARNING **: 11:06:20.519: GVfs metadata is not supported. Fallba
ck to TeplMetadataManager. Either GVfs is not correctly installed or GVfs metadata a
re not supported on this platform. In the latter case, you should configure Tepl wit
h --disable-gvfs-metadata.
wq
^Z
[1]+  Stopped                  sudo gedit /etc/modprobe.d/blacklist.conf
xiaoli@xiaoli-KLVC-WXX9:~$ sudo update-initramfs -u
update-initramfs: Generating /boot/initrd.img-5.15.0-58-generic
I: The initramfs will attempt to resume from /dev/nvme0n1p7
I: (UUID=1e7230ed-440d-4817-9966-af93591f44c5)
I: Set the RESUME variable to override this.

```

1. 重启

2. 进入NVIDIA官网CUDA下载页面<https://developer.nvidia.com/cuda-toolkit-archive>选择
CUDA Toolkit 11.2.0(December 2020)，依次选择
Linux→x86_64→Ubuntu→20.04→runfile(local)
3. 在终端输入sudo apt-get install freeglut3-dev build-essential libx11-dev libxmu-
dev libxi-dev libgl1-mesa-glx libglu1-mesa libglu1-mesa-dev安装依赖库文件
4. 在终端输入wget
https://developer.download.nvidia.com/compute/cuda/11.2.0/local_installers/cuda_11.2.0_460.27.04_linux.run和sudo sh cuda_11.2.0_460.27.04_linux.run安装CUDA。接
下来会弹出两个页面，在第一个页面输入accept、回车，在第二个页面按空格取消Driver勾选，然后点击

Install、等待。

```

+ xiaoli@xiaoli-KLVC-WXX9: ~
xiaoli@xiaoli-KLVC-WXX9:~$ wget https://developer.download.nvidia.com/compute/cuda/11.2.0/local_installers/cuda_11.2.0_460.27.04_linux.run
--2023-04-04 11:48:31-- https://developer.download.nvidia.com/compute/cuda/11.2.0/local_installers/cuda_11.2.0_460.27.04_linux.run
Connecting to 127.0.0.1:7890... connected.
Proxy request sent, awaiting response... 200 OK
Length: 3046790184 (2.8G) [application/octet-stream]
Saving to: 'cuda_11.2.0_460.27.04_linux.run'

cuda_11.2.0_460.27. 100%[=====>] 2.84G 11.5MB/s in 6m 10s

2023-04-04 11:54:42 (7.86 MB/s) - 'cuda_11.2.0_460.27.04_linux.run' saved [3046790184/3046790184]

xiaoli@xiaoli-KLVC-WXX9:~$ sudo sh cuda_11.2.0_460.27.04_linux.run
=====
= Summary =
=====

Driver: Not Selected
Toolkit: Installed in /usr/local/cuda-11.2/
Samples: Installed in /home/xiaoli/

```

5. 配置环境变量

```

$ export PATH=/usr/local/cuda-10.1/bin${PATH:+:${PATH}}
$ export LD_LIBRARY_PATH=/usr/local/cuda-10.1/lib64\
    ${LD_LIBRARY_PATH:+:${LD_LIBRARY_PATH}}

```

6. 在终端输入 `source ~/.bashrc` 使环境变量生效。

7. 查看CUDA安装信息

```

+ xiaoli@xiaoli-KLVC-WXX9: ~/NVIDIA_CUDA-11.2_Samples/1_Uutilities/dev...
xiaoli@xiaoli-KLVC-WXX9:~/NVIDIA_CUDA-11.2_Samples/1_Uutilities/deviceQuery$ nvcc -V
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2020 NVIDIA Corporation
Built on Mon_Nov_30_19:08:53_PST_2020
Cuda compilation tools, release 11.2, V11.2.67
Build cuda_11.2.r11.2/compiler.29373293_0

```

8. CUDA测试。进入NVIDIA CUDA示例包，其位于 `/home/xiaoli/NVIDIA_CUDA-11.2_Samples`，在该文件夹下打开终端，并输入 `make`。然后进入 `1_Uutilities/deviceQuery` 文件夹，并在终端执

行./deviceQuery命令，输出结果result=PASS表示安装成功。

```

+ xiaoli@xiaoli-KLVC-WXX9: ~/NVIDIA_CUDA-11.2_Samples/1...
xiaoli@xiaoli-KLVC-WXX9:~/NVIDIA_CUDA-11.2_Samples/1_Uutilities/deviceQuery$ ./de
viceQuery
./deviceQuery Starting...

  CUDA Device Query (Runtime API) version (CUDA RT static linking)

Detected 1 CUDA Capable device(s)

Device 0: "NVIDIA GeForce MX350"
  CUDA Driver Version / Runtime Version      12.0 / 11.2
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             2001 MBytes (2098331648 bytes)
  ( 5) Multiprocessors, (128) CUDA Cores/MP: 640 CUDA Cores
  GPU Max Clock rate:                        1468 MHz (1.47 GHz)
  Memory Clock rate:                         3504 Mhz
  Memory Bus Width:                          64-bit
  L2 Cache Size:                             524288 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536)
, 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total shared memory per multiprocessor:     98304 bytes
  Total number of registers available per block: 65536
  Warp size:                                  32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
  Run time limit on kernels:                   Yes
  Integrated GPU sharing Host Memory:          No
  Support host page-locked memory mapping:     Yes
  Alignment requirement for Surfaces:          Yes
  Device has ECC support:                      Disabled
  Device supports Unified Addressing (UVA):     Yes
  Device supports Managed Memory:              Yes
  Device supports Compute Preemption:          Yes
  Supports Cooperative Kernel Launch:          Yes
  Supports MultiDevice Co-op Kernel Launch:    Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simu
ltaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.0, CUDA Runtime Vers
ion = 11.2, NumDevs = 1
Result = PASS

```

Lab 2 - 排序算法的并行及优化（验证）

Lab 3 - 矩阵乘法的并行及优化（验证）

Lab 4 - 快速傅里叶变换的并行实现（验证）

Lab 5 - 常用图像处理算法的并行及优化（设计）

Appendix

仅以此记录一下自己被困扰了一天的问题。以下是我写的第一个测试OpenMP的C语言代码，其过程是近似计算PI的值

```
#include <stdio.h>
#include <time.h>
#include <omp.h>

void sum(){
    int sum = 0;
    for(int i = 0; i < 1000000000; i++){
        sum++;
    }
}

void parallel(){
    clock_t start, end;
    start = clock();
    # pragma omp parallel for
    for(int i = 0; i < 100; i++){
        sum();
    }
    end = clock();
    printf("Parallel time: %ld \n", end - start);
}

void no_parallel(){
    clock_t start, end;
    start = clock();
    for(int i = 0; i < 100; i++){
        sum();
    }
    end = clock();
    printf("Serial time: %ld \n", end - start);
}
```

```
int main() {  
    parallel();  
    no_parallel();  
    return 0;  
}
```

```
• xiaoli@xiaoli-KLVC-WXX9:~/Project/Parallel-computing$ gcc compute_pi.c -o output  
• xiaoli@xiaoli-KLVC-WXX9:~/Project/Parallel-computing$ ./output  
Elapsed time: 0.583642 seconds  
pi = 3.141593  
• xiaoli@xiaoli-KLVC-WXX9:~/Project/Parallel-computing$ gcc -fopenmp compute_pi.c -o output  
• xiaoli@xiaoli-KLVC-WXX9:~/Project/Parallel-computing$ ./output  
Elapsed time: 0.600783 seconds  
pi = 3.141593
```

但是输出结果却令我大为震惊，因为开了并行竟然比不开更浪费时间，虽然我一开始以为可能是老师上课说的那种情况——并行的开销比并行的收益更大，但是当我把参数量调大之后发现这个现象仍然存在，于是我上网进行了搜索，终于发现原来是时间的测量方法使用错误。`clock()`记录的是CPU的滴答数，当并行多个进程同时计算，CPU滴答数成倍增加，所以我们得到的差值并不是真实的时间数，OpenMP提供的`omp_get_wtime()`才记录的是真实的运行时间，当我把时间测量函数修改后发现代码运行正常，结果为